

Robert Bielas
laboratorium z Netezzy - raport

UWAGA: proszę zwracać uwagę na komentarze pod screenami/tekstami planera, gdyż pokazują one moją interpretację wyników.

Część pierwsza, zad. 11

Najlepsza dystrybucja tabel dla tabel w LABDB: w tabelach joinowanych (orders i lineitem, part i lineitem, supplier i part sup itp.) zaproponowałbym dystrybucję wg kolumny, która łączy odpowiednie pary tabel. Jeśli nie zależy nam na informacjach z kilku tabel, ale z jednej, zwróciłbym uwagę na to, którą operację wyszukiwania będziemy najczęściej wykonywać: np. wg. daty. Wtedy zdystrybuowałbym tabelę, posortowaną wg interesującej nas kolumny, i zdystrybuował po kluczu którego zbiór unikalnych wartości ma moc ≥ 22 (tyle najwyżej data slice'ów jest w serwerze Netezzy wykorzystywanej na laboratorium).

Część druga, zad. 1

```
select count(*) from paymentreport, rezultat:  
100000000
```

```
select count(*) from paymentreportdetail, rezultat:  
600000000
```

```
select count(*) from company, rezultat: 22
```

```
select count(*) from COMPANYHEADQUARTER, rezultat: 43  
select count(*) from Paymentreport_age  
100000000
```

```
select count(*) from Paymentreport_gender, rezultat:  
100000000
```

Wniosek: tabele są pokaźnych rozmiarów

zad. 2 rezultat:

DATASLICEID	↓ Σ ▾ ▹ ▸	rows	↓ Σ ▾ ▹ ▸
1		4546633	
2		4546634	
3		4546634	
4		4546634	
5		4546634	
6		4545108	
7		4545108	
8		4545108	
9		4545108	
10		4545108	
11		4545108	
12		4545108	
13		4545108	
14		4545108	
15		4545108	
16		4545108	
17		4545108	
18		4545107	
19		4545107	
20		4545107	
21		4545107	
22		4545107	

zad. 3

Dystrybucja tabeli paymentreport_gender (po kluczu gender):

```
select datasliceid, count(datasliceid) as "rows"  
from PAYMENTREPORT_GENDER  
group by datasliceid  
order by 1
```

DATASLICEID	↓ ∑ ▾ ▹ ▸	rows	↓ ∑ ▾ ▹ ▸
10		49999997	
15		50000003	

Dystrybuja tabeli paymentreport_age (po kluczu age)

```
select datasliceid, count(datasliceid) as "rows"  
from PAYMENTREPORT_AGE  
group by datasliceid  
order by 1
```

DATASLICEID	↓ ∑ ▾ ▹ ▸	rows	↓ ∑ ▾ ▹ ▸
1		5000542	
2		5997064	
3		4999551	
4		5003441	
5		4995738	
6		5000843	
7		5000580	
8		5003587	
9		5000312	
10		5002074	
11		4004265	
12		3997798	
13		3999722	
14		4001385	
15		3998297	
16		3995568	
17		4003395	
18		4002063	
19		4999312	
20		3998698	
21		3998174	
22		3997591	

Po utworzeniu tabel w mojej bazie na tych dystrybucjach rozmieszczenie krotek było identyczne.

zad. 4

```
137 select gender, avg(salary)
138 from PAYMENTREPORT
139 group by gender
140
```

<

Output Result 1

Drag a column header here to group by that column.

GENDER	AVG
M	26003,028533
F	25997,030155

U_bielas 1 query, 1,300 sec, 2 rows laby.sql

```
141 select age, avg(salary)
142 from PAYMENTREPORT
143 group by age
144 order by 1
```

<

Output Result 1

Drag a column header here to group by that column.

AGE	AVG
0	26003,247861
1	26001,350730

U_bielas 1 query, 1,356 sec, 68 rows * laby.sql

Sprawdzenie czasu wykonania obu agregacji dla dystrybucji po age:

```
147 select age, avg(salary)
148 from PAYMENTREPORT_age
149 group by age
150 order by 1
151
```

<

Output Result 1

Drag a column header here to group by that column.

AGE	AVG
0	26003,247861
1	26001,350730
2	26005,294429

U_bielas 1 query, 1,350 sec, 68 rows laby.sql

```

157 select gender, avg(salary)
158 from PAYMENTREPORT_age
159 group by gender
160

```

GENDER	AVG
F	25997,030155
M	26003,028533

U_bielas 1 query, 1,425 sec, 2 rows * laby.sql

Czas wykonania jest rozsądny, gdyż nie występuje data skew - dane wydystribuowane po age są równomiernie rozłożone po większej ilości SPU - jak w przypadku dystrybucji po id. Sprawdzenie czasu wykonania obu agregacji dla dystrybucji po gender:

```

152 select age, avg(salary)
153 from PAYMENTREPORT_gender
154 group by age
155 order by 1

```

AGE	AVG
0	26003,247861
1	26001,350730
2	26005,294429
3	26010,007374

U_bielas 1 query, 9,016 sec, 68 rows * laby.sql

```

161 select gender, avg(salary)
162 from PAYMENTREPORT_gender
163 group by gender
164
165

```

GENDER	AVG
M	26003,028533
F	25997,030155

U_bielas 1 query, 9,043 sec, 2 rows * laby.sql

Tutaj nie dziwi nas, że czasy są podobne - w przetwarzanie są zaangażowane tylko dwa S-Blade'y (bo do dwóch wydystribuowano dane) - reszta śpi.

zad. 5

```
select count(*) from PAYMENTREPORTDETAIL_ALL_BONUSCODE
1600000000
```

```
select count(*) from PAYMENTREPORTDETAIL_ALL2
3200000000
```

```
select count(*) from PAYMENTREPORTDETAIL_ALL2_BONUSCODE
3200000000
```

```
select count(*) from PAYMENTREPORTDETAIL_ALL
1600000000
```

dystrybucja w paymentreportdetail_all

DATASLICEID	rows
1	72746141
2	72746142
3	72746142
4	72746142
5	72746142
6	72721723
7	72721723
8	72721723
9	72721723
10	72721723
11	72721723
12	72721723
13	72721723
14	72721723
15	72721723
16	72721723
17	72721723
18	72721723
19	72721723
20	72721723
21	72721723
22	72721723

dystrybucja w paymentreportdetail_all_bonuscode

```
151 select datasliceid, count(datasliceid) as "rows"
152 from PAYMENTREPORTDETAIL_ALL_BONUSCODE
153 group by DATASLICEID
154 order by 1
155
156
```

Output Result 1

Drag a column header here to group by that column.

DATASLICEID	rows
1	400006187
2	800027467
3	399966346

dystrybucja w paymentreportdetail_all2

```
155
156 |select datasliceid, count(datasliceid) as "rows"
157 |from PAYMENTREPORTDETAIL_ALL2
158 |group by DATASLICEID
159 |order by 1
```

Output Result 1

Drag a column header here to group by that column.

DATASLICEID	rows
1	145492284
2	145492285
3	145492285
4	145492285
5	145492285
6	145443446
7	145443446
8	145443446
9	145443446
10	145443446
11	145443446
12	145443446
13	145443446
14	145443446
15	145443446
16	145443446
17	145443445
18	145443445
19	145443445
20	145443445
21	145443445
22	145443445

dystrybucja w paymentreportdetail_all2_bonuscode:

```
161
162 |select datasliceid, count(datasliceid) as "rows"
163 |from PAYMENTREPORTDETAIL_ALL2_BONUSCODE
164 |group by DATASLICEID
165 |order by 1
166 |
```

Output Result 1

Drag a column header here to group by that column.

DATASLICEID	rows
1	800037916
2	1599976875
3	799985209

tabela paymentreportdetail_all jest zdystrybuowana na id
tabela paymentreportdetail_all2 jest rozdystrybuowana na id
tabela paymentreportdetail_all2_bonuscode jest rozdystrybuowana na bonuscode
tabela paymentreportdetail_all_bonuscode jest rozdystrybuowana na bonuscode

Polecenia agregujące:

```
select date, sum(bonus)
from PAYMENTREPORTDETAIL_ALL
group by date
```

czas: 13.916 sec

```
select bonuscode, sum(bonus)
from PAYMENTREPORTDETAIL_REPORTID
group by bonuscode
```

czas: 4.793 sec

```
select bonuscode, sum(bonus)
from PAYMENTREPORTDETAIL_ALL_BONUSCODE
group by bonuscode
czas: 85.993 sec (!!!)
```

Ostatnia agregacja wykonywała się najwolniej, ponieważ rozdystrybuowano ją na bonuscode(sa 3). Mamy tu do czynienia ze zjawiskiem data_skew: tylko 3 data slice'y były zaangażowane w przetwarzanie (do trzech slice'ów zostały zdystrybuowane dane).

Druga agregacja jest najszybsza, gdyż dane zostały przechowane równomiernie w największej ilości slice'ów. Więcej S-Bladów przetwarzało mniej danych -> wynik marzenie.

6.

Czas wykonania query dla paymentreportdetail_reportid2 (dystrybucja po payment_report, posortowane po dacie):

The screenshot shows a SQL query editor with the following query:

```
184 select bonuscode, sum(bonus)
185 from paymentreportdetail_reportid2
186 where date = '2015-10-20'
187 group by bonuscode
188
```

Below the query, there is a table with the results:

BONUSCODE	SUM
1	33866578235,00
0	16908522894,00
2	16926967834,00

At the bottom, it shows the execution time: 1 query, 0,264 sec, 3 rows * laby.sql

DATASLICEID	↓ Σ ∇ ▢	rep2	↓ Σ ∇ ▢
1		27279798	
2		27279804	
3		27279804	
4		27279804	
5		27279804	
6		27270648	
7		27270648	
8		27270648	
9		27270648	
10		27270648	
11		27270648	
12		27270648	
13		27270648	
14		27270648	
15		27270648	
16		27270648	
17		27270648	
18		27270642	
19		27270642	
20		27270642	
21		27270642	
22		27270642	

Czas wykonania query dla paymentreportdetail_reporttid

```

190 select bonuscode, sum(bonus)
191 from paymentreportdetail_reporttid
192 where date = '2015-10-20'
193 group by bonuscode

```

Output Result 1

Drag a column header here to group by that column.

BONUSCODE	SUM
1	33866578235,00
2	16926967834,00
0	16908522894,00

U_bielas 1 query, 5,313 sec., 3 rows laby.sql

DATASLICEID	rep
1	27279798
2	27279804
3	27279804
4	27279804
5	27279804
6	27270648
7	27270648
8	27270648
9	27270648
10	27270648
11	27270648
12	27270648
13	27270648
14	27270648
15	27270648
16	27270648
17	27270648
18	27270642
19	27270642
20	27270642
21	27270642
22	27270642

Czas wykonania query dla paymentreportdetail_date:

```

195 select bonuscode, sum(bonus)
196 from paymentreportdetail_date
197 where date = '2015-10-20'
198 group by bonuscode

```

Output Result 1

Drag a column header here to group by that column.

BONUSCODE	SUM
2	16926967834,00
1	33866578235,00
0	16908522894,00

U_bielas 1 query, 2,619 sec., 3 rows * laby.sql

DATASLICEID	date
1	30000824
2	29999864
3	30011297
4	29998409
5	29993986
6	29999312
7	24997732
8	20003703
9	20002352
10	20000642
11	20005831
12	20001057
13	24997217
14	30001279
15	30000847
16	29997767
17	29999905
18	30004882
19	29990426
20	30003865
21	29996971
22	29991832

Najbardziej efektywnie zostało wykonane query dla dystrybucji po payment_report, gdzie dane zostały posortowane po dacie (czyli w tabeli paymentreportdetail_reporttid2).

Czasy wykonania dla większych tabel:

Czas wykonania dla paymentreportdetail_all:

```

200 select bonuscode, sum(bonus)
201 from paymentreportdetail_all
202 where date = '2015-10-20'
203 group by bonuscode

```

Output Result 1

Drag a column header here to group by that column.

BONUSCODE	SUM
1	130523816882,00
2	65283759767,00
0	65220461642,00

U_bielas 1 query, 13,501 sec., 3 rows * laby.sql

Czas wykonania dla paymentreportdetail_all2:

```
205 select bonuscode, sum(bonus)
206 from paymentreportdetail_all2
207 where date = '2015-10-20'
208 group by bonuscode
```

Output Result 1

Drag a column header here to group by that column.

BONUSCODE	SUM
1	394162785331,00
0	197222832268,00
2	197158522225,00

U_bielas 1 query, 26,751 sec, 3 rows * laby.sql

Czas wykonania jest proporcjonalny do ilości krotek w tabeli (w _all jest ich dwa razy mniej niż w _all2 -> w _all czas wykonania był 2 razy mniejszy). Pamiętajmy, że dane w tych tabelach zostały rozdystrybuowane na tym samym kluczu.

7.

```
210 select company_headquarter, sum(bonus)
211 from paymentreportdetail as pd join PAYMENTREPORT as p
212 on pd.payment_report = p.id
213 group by company_headquarter
```

Output Result 1

Drag a column header here to group by that column.

COMPANY_HEADQUARTER	SUM
3	94343850807,00
22	94554352576,00
27	94429664253,00
38	94401200179,00

U_bielas 1 query, 29,879 sec, 44 rows laby.sql

QUERY VERBOSE PLAN:

Node 1.

[SPU Sequential Scan table "PAYMENTREPORTDETAIL" as "PD" {(PD."ID")}]

-- Estimated Rows = 599005119, Width = 8, Cost = 0.0 .. 11718.8, Conf = 90.0 (FACT)

Restrictions:

(PD.PAYMENT_REPORT NOTNULL)

Projections:

1:PD.BONUS 2:PD.PAYMENT_REPORT

Cardinality:

PD.PAYMENT_REPORT 100.0M (JIT)

[SPU Distribute on {(PD.PAYMENT_REPORT)}]

[HashIt for Join]

Node 2.

[SPU Sequential Scan table "PAYMENTREPORT" as "P" {(P."ID")}]

-- Estimated Rows = 100000000, Width = 8, Cost = 0.0 .. 3595.5, Conf = 80.0 (FACT)

Restrictions:

(P."ID" NOTNULL)

Projections:

1:P.COMPANY_HEADQUARTER 2:P."ID"

Node 3.

[SPU Hash Join Stream "Node 2" with Temp "Node 1"
{(PD.PAYMENT_REPORT,P."ID")}]

-- Estimated Rows = 599005119, Width = 8, Cost = 39316.9 .. 84826.1, Conf = 64.0

Restrictions:

(P."ID" = PD.PAYMENT_REPORT)

Projections:

1:P.COMPANY_HEADQUARTER 2:PD.BONUS

Cardinality:

PD.PAYMENT_REPORT 85.2M (Adjusted)

Node 4.

[SPU Group]

-- Estimated Rows = 44, Width = 8, Cost = 39316.9 .. 98439.9, Conf = 0.0

Projections:

1:P.COMPANY_HEADQUARTER 2:PD.BONUS

[SPU Return]

[HOST Merge Group]

Node 5.

[Host Aggregate]

-- Estimated Rows = 44, Width = 20, Cost = 39316.9 .. 98439.9, Conf = 0.0

Projections:

1:P.COMPANY_HEADQUARTER 2:SUM(PD.BONUS)

[Host Return]

Z planu wynika, że jest wykonywana redystrybucja po kluczu paymentreport. Mamy do czynienia z single redistribution. Ponieważ taka sytuacja zachodzi, konieczność owej redystrybucji (wyciąganie odpowiednich danych z data slice'ów i przesłanie ich do odpowiedniego S-Blade'a) wpływa negatywnie na czas wykonania.

215	select company_headquarter, sum(bonus)
216	from paymentreportdetail_reportid as pd join PAYMENTREPORT as p
217	on pd.payment_report = p.id
218	group by company_headquarter
219	

Output
Result 1

Drag a column header here to group by that column.

COMPANY_HEADQUARTER	SUM
14	94344254872.00
29	94346100735.00
21	94253911956.00
28	94360543150.00

Pos: 5241 (row: 218, col: 29) U_bielas 1 query, 14,634 sec., 44 rows * laby.sql

QUERY VERBOSE PLAN:

Node 1.

[SPU Sequential Scan table "PAYMENTREPORT" as "P" {(P."ID")}]

-- Estimated Rows = 100000000, Width = 8, Cost = 0.0 .. 3595.5, Conf = 80.0 (FACT)

Restrictions:

(P."ID" NOTNULL)

Projections:

1:P.COMPANY_HEADQUARTER 2:P."ID"

[HashIt for Join]

Node 2.

[SPU Sequential Scan table "PAYMENTREPORTDETAIL_REPORTID" as "PD" {(PD.PAYMENT_REPORT)}]

-- Estimated Rows = 599426939, Width = 8, Cost = 0.0 .. 11718.8, Conf = 90.0 (FACT)

Restrictions:

(PD.PAYMENT_REPORT NOTNULL)

Projections:

1:PD.BONUS 2:PD.PAYMENT_REPORT

Cardinality:

PD.PAYMENT_REPORT 100.0M (JIT)

Node 3.

[SPU Hash Join Stream "Node 2" with Temp "Node 1" {(P."ID",PD.PAYMENT_REPORT)}]

-- Estimated Rows = 599426939, Width = 8, Cost = 3595.5 .. 72817.7, Conf = 64.0

Restrictions:

(PD.PAYMENT_REPORT = P."ID")

Projections:

1:P.COMPANY_HEADQUARTER 2:PD.BONUS

Cardinality:

PD.PAYMENT_REPORT 85.2M (Adjusted)

Node 4.

[SPU Group]

-- Estimated Rows = 44, Width = 8, Cost = 3595.5 .. 86441.0, Conf = 0.0

Projections:

1:P.COMPANY_HEADQUARTER 2:PD.BONUS

[SPU Return]

[HOST Merge Group]

Node 5.

[Host Aggregate]

-- Estimated Rows = 44, Width = 20, Cost = 3595.5 .. 86441.0, Conf = 0.0

Projections:

1:P.COMPANY_HEADQUARTER 2:SUM(PD.BONUS)

[Host Return]

Czas wykonywania jest krótszy, gdyż nie trzeba dokonywać redystrybucji (jak wynika z planera - nie ma wzmianki distribute on(...)).

8.

```
219
220 select c.company_name, sum(bonus)
221 from paymentreportdetail_reportid as pd join paymentreport as p
222 on pd.payment_report = p.id join companyheadquarter as ch
223 on ch.id = p.company_headquarter join company as c
224 on c.id = ch.company
225 group by c.company_name
```

<

Output Result 1

Drag a column header here to group by that column.

COMPANY_NAME	SUM
Yozio	283449163204,00
Feedbug	283120922089,00
Livetube	188673549830,00
Dynabox	236186082003,00
Eayo	188850285736,00
Thoughtstorm	188918712105,00
Wikizz	94554352576,00
Browseblab	283041012034,00

Pos: 5530 (row: 225, col: 24) U_bielas 1 query, 19,497 sec., 21 rows laby.sql

QUERY VERBOSE PLAN:

Node 1.

[SPU Sequential Scan table "PAYMENTREPORTDETAIL_REPORTID" as "PD" {(PD.PAYMENT_REPORT)}]

-- Estimated Rows = 599426939, Width = 8, Cost = 0.0 .. 11718.8, Conf = 90.0 (FACT)

Restrictions:

(PD.PAYMENT_REPORT NOTNULL)

Projections:

1:PD.BONUS 2:PD.PAYMENT_REPORT

Cardinality:

PD.PAYMENT_REPORT 100.0M (JIT)

[HashIt for Join]

Node 2.

[SPU Sequential Scan table "COMPANYHEADQUARTER" as "CH" {(CH."ID")}]

-- Estimated Rows = 43, Width = 8, Cost = 0.0 .. 0.0, Conf = 100.0

Restrictions:

((CH."ID" NOTNULL) AND (CH.COMPANY NOTNULL))

Projections:

1:CH."ID" 2:CH.COMPANY

[HashIt for Join]

Node 3.

[SPU Sequential Scan table "COMPANY" as "C" {(C."ID")}]

-- Estimated Rows = 22, Width = 54, Cost = 0.0 .. 0.0, Conf = 80.0

Restrictions:

(C."ID" NOTNULL)

Projections:

1:C.COMPANY_NAME 2:C."ID"

Node 4.

[SPU Hash Join Stream "Node 3" with Temp "Node 2" {(C."ID")}]

-- Estimated Rows = 22, Width = 54, Cost = 0.0 .. 0.0, Conf = 64.0

Restrictions:

(C."ID" = CH.COMPANY)

Projections:

1:CH."ID" 2:C.COMPANY_NAME

Cardinality:

CH."ID" 22 (Adjusted)

CH.COMPANY 22 (Adjusted)

[SPU Broadcast]

[HashIt for Join]

Node 5.

[SPU Sequential Scan table "PAYMENTREPORT" as "P" {(P."ID")}]

-- Estimated Rows = 99935252, Width = 8, Cost = 0.0 .. 1139.2, Conf = 90.0 [BT: MaxPages=890 TotalPages=19580] (JIT-Stats)

Restrictions:

((P."ID" NOTNULL) AND (P.COMPANY_HEADQUARTER NOTNULL))

Projections:

1:P."ID" 2:P.COMPANY_HEADQUARTER

Cardinality:

P.COMPANY_HEADQUARTER 43 (Adjusted)

Node 6.

[SPU Hash Join Stream "Node 5" with Temp "Node 4" {(P."ID")}]

-- Estimated Rows = 51129664, Width = 54, Cost = 0.0 .. 10451.4, Conf = 51.2

Restrictions:

(P.COMPANY_HEADQUARTER = CH."ID")

Projections:

1:C.COMPANY_NAME 2:P."ID"

Cardinality:

CH."ID" 22 (Adjusted)

CH.COMPANY 22 (Adjusted)

P.COMPANY_HEADQUARTER 22 (Adjusted)

Node 7.

[SPU Hash Join Stream "Node 6" with Temp "Node 1" {(PD.PAYMENT_REPORT,P."ID")}]

-- Estimated Rows = 302930461, Width = 54, Cost = 36343.7 .. 76782.2, Conf = 46.1

Restrictions:

(P."ID" = PD.PAYMENT_REPORT)

Projections:

1:C.COMPANY_NAME 2:PD.BONUS

Cardinality:

CH.COMPANY 22 (Adjusted)

Node 8.

[SPU Group]

-- Estimated Rows = 22, Width = 54, Cost = 36343.7 .. 83667.0, Conf = 0.0

Projections:

1:C.COMPANY_NAME 2:PD.BONUS

[SPU Return]

[HOST Merge Group]

Node 9.

[Host Aggregate]

-- Estimated Rows = 22, Width = 66, Cost = 36343.7 .. 83667.0, Conf = 0.0

Projections:

1:C.COMPANY_NAME 2:SUM(PD.BONUS)

[Host Return]

227	select c.company_name, sum(bonus)
228	from paymentreportdetail as pd join paymentreport as p
229	on pd.payment_report = p.id join companyheadquarter as ch
230	on ch.id = p.company_headquarter join company as c
231	on c.id = ch.company
232	group by c.company_name
233	

Output		Result 1			
Drag a column header here to group by that column.					
COMPANY_NAME	▼	SUM	↓	Σ	▼
Yozio		283449163204,00			
Browseblab		283041012034,00			
Dynabox		236186082003,00			
Shuffletag		283224968548,00			
Eayo		188850285736,00			
Youspan		283162689779,00			
Feedbug		283120922089,00			
Realpoint		94420412810,00			

Pos: 5794 (row: 232, col: 24)	U_bielas	1 query, 31,465 sec, 21 rows * laby.sql
-------------------------------	----------	---

QUERY VERBOSE PLAN:

Node 1.
[SPU Sequential Scan table "PAYMENTREPORTDETAIL" as "PD" {(PD."ID")}]
-- Estimated Rows = 599005119, Width = 8, Cost = 0.0 .. 11718.8, Conf = 90.0 (FACT)
Restrictions:
(PD.PAYMENT_REPORT NOTNULL)
Projections:
1:PD.BONUS 2:PD.PAYMENT_REPORT
Cardinality:
PD.PAYMENT_REPORT 100.0M (JIT)
[SPU Distribute on {(PD.PAYMENT_REPORT)}]
[HashIt for Join]
Node 2.
[SPU Sequential Scan table "COMPANYHEADQUARTER" as "CH" {(CH."ID")}]
-- Estimated Rows = 43, Width = 8, Cost = 0.0 .. 0.0, Conf = 100.0
Restrictions:
((CH."ID" NOTNULL) AND (CH.COMPANY NOTNULL))
Projections:
1:CH."ID" 2:CH.COMPANY
[HashIt for Join]
Node 3.
[SPU Sequential Scan table "COMPANY" as "C" {(C."ID")}]
-- Estimated Rows = 22, Width = 54, Cost = 0.0 .. 0.0, Conf = 80.0

Restrictions:

(C."ID" NOTNULL)

Projections:

1:C.COMPANY_NAME 2:C."ID"

Node 4.

[SPU Hash Join Stream "Node 3" with Temp "Node 2" {(C."ID")}]

-- Estimated Rows = 22, Width = 54, Cost = 0.0 .. 0.0, Conf = 64.0

Restrictions:

(C."ID" = CH.COMPANY)

Projections:

1:CH."ID" 2:C.COMPANY_NAME

Cardinality:

CH."ID" 22 (Adjusted)

CH.COMPANY 22 (Adjusted)

[SPU Broadcast]

[HashIt for Join]

Node 5.

[SPU Sequential Scan table "PAYMENTREPORT" as "P" {(P."ID")}]

-- Estimated Rows = 99935252, Width = 8, Cost = 0.0 .. 1139.2, Conf = 90.0 [BT: MaxPages=890
TotalPages=19580] (JIT-Stats)

Restrictions:

((P."ID" NOTNULL) AND (P.COMPANY_HEADQUARTER NOTNULL))

Projections:

1:P."ID" 2:P.COMPANY_HEADQUARTER

Cardinality:

P.COMPANY_HEADQUARTER 43 (Adjusted)

Node 6.

[SPU Hash Join Stream "Node 5" with Temp "Node 4" {(P."ID")}]

-- Estimated Rows = 51129664, Width = 54, Cost = 0.0 .. 10451.4, Conf = 51.2

Restrictions:

(P.COMPANY_HEADQUARTER = CH."ID")

Projections:

1:C.COMPANY_NAME 2:P."ID"

Cardinality:

CH."ID" 22 (Adjusted)

CH.COMPANY 22 (Adjusted)

P.COMPANY_HEADQUARTER 22 (Adjusted)

Node 7.

[SPU Hash Join Stream "Node 6" with Temp "Node 1" {(PD.PAYMENT_REPORT,P."ID")}]

-- Estimated Rows = 302717287, Width = 54, Cost = 36327.2 .. 86000.2, Conf = 46.1

Restrictions:

(P."ID" = PD.PAYMENT_REPORT)

Projections:

1:C.COMPANY_NAME 2:PD.BONUS

Cardinality:

CH.COMPANY 22 (Adjusted)

Node 8.

[SPU Group]

-- Estimated Rows = 22, Width = 54, Cost = 36327.2 .. 92880.1, Conf = 0.0

Projections:

1:C.COMPANY_NAME 2:PD.BONUS

[SPU Return]

[HOST Merge Group]

Node 9.

[Host Aggregate]

-- Estimated Rows = 22, Width = 66, Cost = 36327.2 .. 92880.1, Conf = 0.0

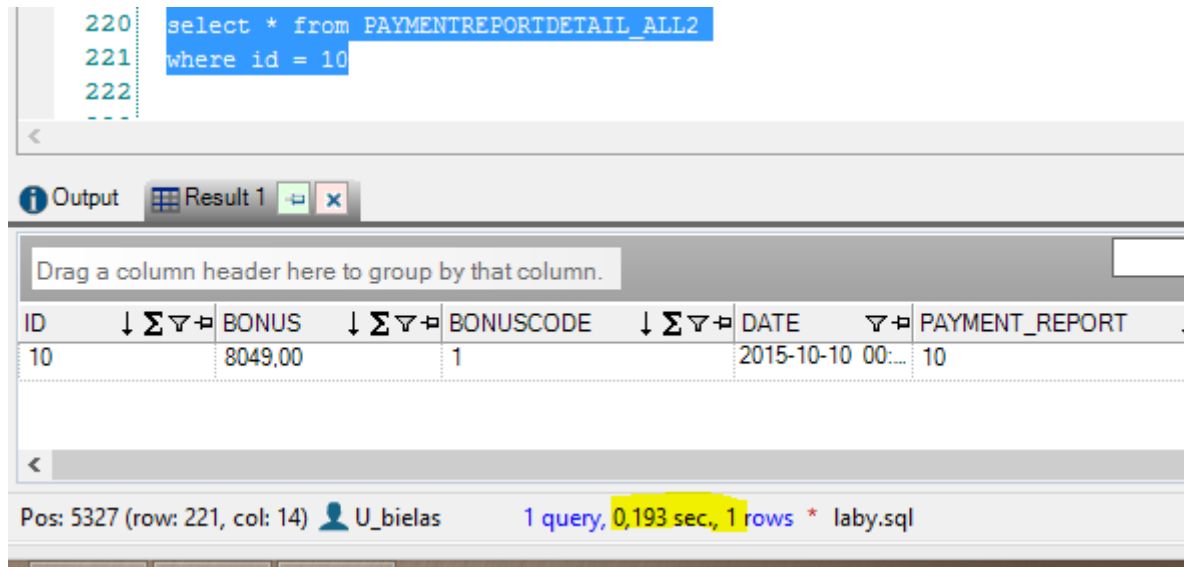
Projections:

1:C.COMPANY_NAME 2:SUM(PD.BONUS)

[Host Return]

W obu powyższych sytuacjach pojawia się wzorzec, o którym wspomniano w zad.7:
tam gdzie jest konieczność zastosowania redystrybucji, negatywnie (czasowo) odbija się
to na wydajności.

9. ciekawostka:



The screenshot shows a SQL query execution interface. At the top, a query is entered in a text area: `select * from PAYMENTREPORTDETAIL_ALL2 where id = 10`. Below the query, there is a tab labeled "Output" and "Result 1". The results are displayed in a table with the following columns: ID, BONUS, BONUSCODE, DATE, and PAYMENT_REPORT. The table contains one row with the following values: ID: 10, BONUS: 8049,00, BONUSCODE: 1, DATE: 2015-10-10 00:..., PAYMENT_REPORT: 10. At the bottom of the interface, there is a status bar that reads: "Pos: 5327 (row: 221, col: 14) U_bielas 1 query, 0,193 sec, 1 rows * laby.sql".

ID	BONUS	BONUSCODE	DATE	PAYMENT_REPORT
10	8049,00	1	2015-10-10 00:...	10

QUERY VERBOSE PLAN:

Node 1.

[SPU Sequential Scan table "PAYMENTREPORTDETAIL_ALL2"
{(PAYMENTREPORTDETAIL_ALL2."ID")}]

-- Estimated Rows = 32000000, Width = 24, Cost = 0.0 .. 68181.8, Conf = 80.0

Restrictions:

(PAYMENTREPORTDETAIL_ALL2."ID" = 10)

Projections:

1:PAYMENTREPORTDETAIL_ALL2."ID" 2:PAYMENTREPORTDETAIL_ALL2.BONUS

3:PAYMENTREPORTDETAIL_ALL2.BONUSCODE

4:PAYMENTREPORTDETAIL_ALL2.DATE

5:PAYMENTREPORTDETAIL_ALL2.PAYMENT_REPORT

[SPU Return]

[Host Return]

223	select * from PAYMENTREPORTDETAIL_ALL2
224	where id between 1000000 and 1000050
225	

Output

Result 1

Drag a column header here to group by that column.

ID	↓ Σ ▾ ▹ ▸	BONUS	↓ Σ ▾ ▹ ▸	BONUSCODE	↓ Σ ▾ ▹ ▸	DATE	▾ ▹ ▸	PAYMENT_F
1000015		6525,00		1		2015-11-06 00:...		1000015
1000037		4309,00		1		2015-11-06 00:...		1000037
1000021		7596,00		0		2015-10-12 00:...		1000021
1000043		2527,00		2		2015-10-10 00:...		1000043
1000047		7286,00		2		2015-10-06 00:...		1000047
1000003		8644,00		2		2015-11-20 00:...		1000003
1000025		5035,00		2		2015-10-13 00:...		1000025

Pos: 5422 (row: 225, col: 1)

U_bielas

1 query, 0,339 sec, 51 rows * laby.sql

QUERY VERBOSE PLAN:

Node 1.

[SPU Sequential Scan table "PAYMENTREPORTDETAIL_ALL2"
{(PAYMENTREPORTDETAIL_ALL2."ID")}]

-- Estimated Rows = 51, Width = 24, Cost = 0.0 .. 68181.8, Conf = 64.0

Restrictions:

((PAYMENTREPORTDETAIL_ALL2."ID" <= 1000050) AND
(PAYMENTREPORTDETAIL_ALL2."ID" >= 1000000))

Projections:

1:PAYMENTREPORTDETAIL_ALL2."ID" 2:PAYMENTREPORTDETAIL_ALL2.BONUS

3:PAYMENTREPORTDETAIL_ALL2.BONUSCODE

4:PAYMENTREPORTDETAIL_ALL2.DATE

5:PAYMENTREPORTDETAIL_ALL2.PAYMENT_REPORT

[SPU Return]

[Host Return]

Czasy wyszukiwania jednego rekordu są porównywalne z czasem wyszukiwania 50 rekordów.
Jest to prawdopodobnie zasługa zastosowanych zone map.